

06 Implement SGD

June 6, 2019

```
In [1]: import warnings
        warnings.filterwarnings("ignore")
        from sklearn.datasets import load_boston
        from random import seed
        from random import randrange
        from csv import reader
        from math import sqrt
        from sklearn import preprocessing
        import numpy as np
        import matplotlib.pyplot as plt
        from sklearn.linear_model import SGDRegressor
        from sklearn import preprocessing
        from sklearn.metrics import mean_squared_error
        import time
```

```
In [27]: X = load_boston().data
        Y = load_boston().target
```

```
In [30]: print(X.shape)
        # 506 samples and 13 features
```

(506, 13)

```
In [31]: print("Mean : {}".format(X.mean()))
        print("Max  : {}".format(X.max()))
        print("Min  : {}".format(X.min()))
```

```
Mean : 70.07396704469443
Max  : 711.0
Min  : 0.0
```

```
In [33]: print(Y.shape)
        # Target is the house price in each sample
```

(506,)

```
In [34]: # scaling the data for faster gradient descent
        scaler = preprocessing.StandardScaler().fit(X)
        X = scaler.transform(X)
```

```
In [35]: print("Mean : {}".format(X.mean()))
        print("Max : {}".format(X.max()))
        print("Min : {}".format(X.min()))
```

```
Mean : -1.1147462804871136e-15
Max : 9.933930601860268
Min : -3.9071933049810337
```

```
In [171]: clf = SGDRegressor()
          clf.fit(X, Y)
          print(mean_squared_error(Y, clf.predict(X)))
```

```
23.028034816777073
```

Implementing our own gradient descent function

```
In [37]: samples = X.shape[0]
        iterations = 300
```

```
In [38]: # Adding the 'bias' column with the features columns
        X = np.c_[np.ones(samples), X]
```

```
In [39]: # shuffle data X and Y
        # https://stackoverflow.com/questions/4601373/better-way-to-shuffle-two-numpy-arrays-
        randomize = np.arange(len(X))
        np.random.shuffle(randomize)
        X = X[randomize]
        Y = Y[randomize]
```

```
In [40]: Y = Y.reshape(Y.shape[0], 1)
```

```
In [41]: X.shape
```

```
Out[41]: (506, 14)
```

```
In [13]: def gradient_descent(X, Y, learning_rate=0.01, iterations=400, batch_Size=0):
        # X is matrix used for predictions along with the added bias column
        # Y target variable
        # learning rate determines the step size in the direction of gradient
        # iterations are the num of times for the algo to run
        # batch_Size=1 means SGD, 0 means GD and specify batch for mini-batch

        # sanity check
        assert(Y.shape == (X.shape[0], 1))
```

```

assert(batch_Size >= 0)
samples = X.shape[0]

start = time.time()

# theta initialised with random
theta = np.random.randn(X.shape[1], 1)
# storing the cost values per iteration to plot later
cost_hist = np.zeros(iterations)

for i_num in range(iterations):
    if batch_Size == 0:
        prediction = np.dot(X,theta)
        theta -= (1/samples) * learning_rate * np.dot(X.T, prediction-Y)
        prediction = np.dot(X,theta)
        cost_hist[i_num] = np.sum(np.square(prediction-Y)) / (2*samples)

    elif batch_Size == 1:
        cost = 0.0
        for s in range(samples):
            # iterate over each of the samples as they are already shuffled
            X_i = X[s,:].reshape(1,X.shape[1])
            Y_i = Y[s].reshape(1,1)
            prediction = np.dot(X_i,theta)
            theta -= learning_rate * np.dot(X_i.T, prediction-Y_i)
            prediction = np.dot(X_i,theta)
            cost += np.sum(np.square(prediction-Y_i)) / 2
        cost_hist[i_num] = cost/samples

    else:
        cost = 0.0
        n_batches = samples//batch_Size
        batches = np.random.permutation(samples - batch_Size +1)[:n_batches]
        for bat in batches:
            # iterate over each of the batches
            X_i = X[bat: bat + batch_Size]
            Y_i = Y[bat: bat + batch_Size]
            prediction = np.dot(X_i,theta)
            theta -= (1/batch_Size)*learning_rate*np.dot(X_i.T,prediction-Y_i)
            prediction = np.dot(X_i,theta)
            cost += np.sum(np.square(prediction-Y_i)) / (2*batch_Size)
        cost_hist[i_num] = cost/batch_Size

end = time.time()
print("Run time : {} ms".format(round((end-start)*1000, 2)))
return theta, cost_hist

```

In [42]: *# Batch gradient descent*

```
theta_gd, cost_hist_gd = gradient_descent(X, Y,
                                           iterations=iterations, batch_Size=0)
```

Run time : 22.72 ms

```
In [43]: # stochastic gradient descent
theta_sgd, cost_hist_sgd = gradient_descent(X, Y,
                                           iterations=iterations, batch_Size=1)
```

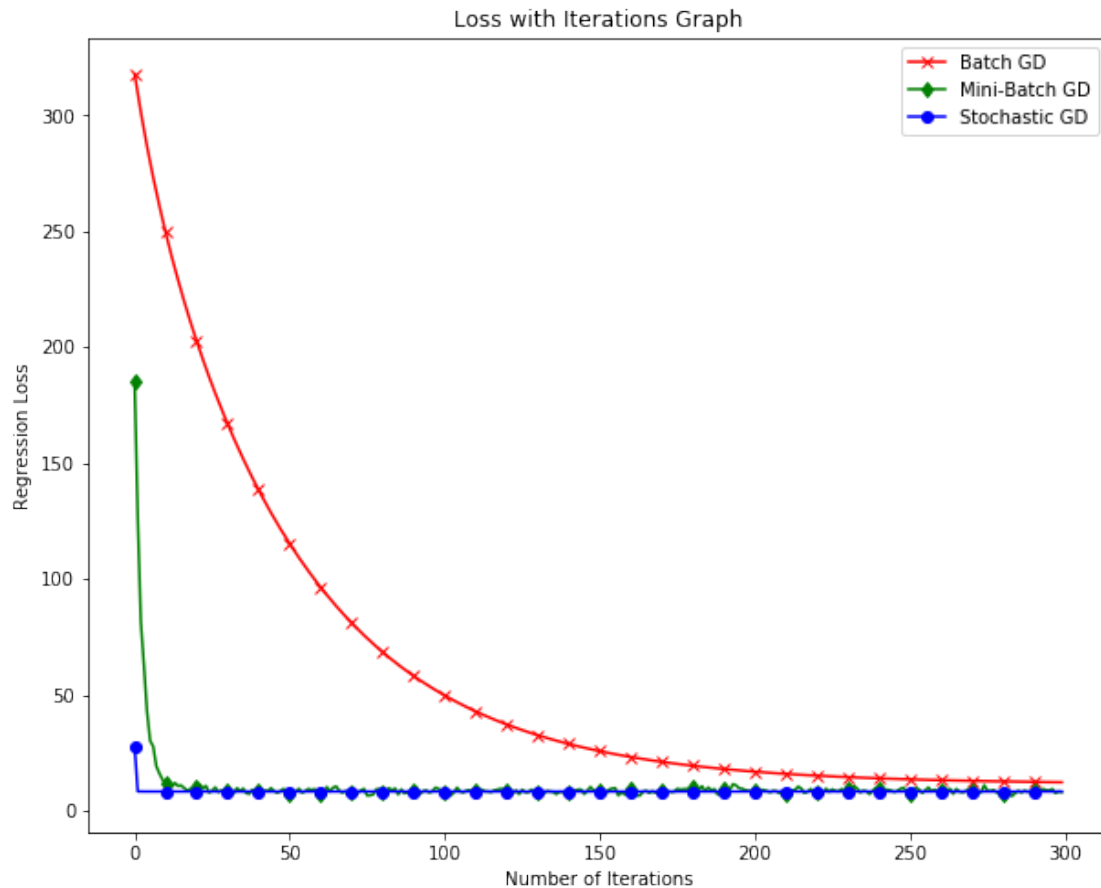
Run time : 2072.1 ms

```
In [44]: # mini-batch gradient descent
theta_mbgd, cost_hist_mbgd = gradient_descent(X, Y,
                                           iterations=iterations, batch_Size=25)
```

Run time : 118.32 ms

Comparing the cost history graphs

```
In [45]: # Plotting the cost values with iterations graph
x_axis = range(iterations)
plt.figure(figsize=(10.0, 8.0))
markers_on = [i for i in range(iterations) if i%10 ==0]
plt.plot(x_axis, cost_hist_gd, '-rx', markevery=markers_on,
         label='Batch GD', )
plt.plot(x_axis, cost_hist_mbgd, '-gd', markevery=markers_on,
         label='Mini-Batch GD')
plt.plot(x_axis, cost_hist_sgd, '-bo', markevery=markers_on,
         label='Stochastic GD')
plt.xlabel("Number of Iterations")
plt.ylabel("Regression Loss")
plt.title("Loss with Iterations Graph")
plt.legend()
plt.show()
```



```
In [46]: clf = SGDRegressor()
         clf.fit(X, Y)
         pred_sklearnSGD = clf.predict(X)
         print(mean_squared_error(Y, pred_sklearnSGD))
```

22.82971463980024

Finding most optimal theta

```
In [47]: # Since the number of features are not too large and our computational
         # power is sufficient, lets directly find the optimal value of theta by
         # using the vectorised normal equation
```

```
theta_best = np.linalg.inv(X.T.dot(X)).dot(X.T).dot(Y)
print(theta_best)
```

```
[[ 2.25328063e+01]
 [-9.28146064e-01]]
```

```
[ 1.08156863e+00]
[ 1.40899997e-01]
[ 6.81739725e-01]
[-2.05671827e+00]
[ 2.67423017e+00]
[ 1.94660717e-02]
[-3.10404426e+00]
[ 2.66221764e+00]
[-2.07678168e+00]
[-2.06060666e+00]
[ 8.49268418e-01]
[-3.74362713e+00]]
```

In [48]: *# Lets compare the Cost value from our gradient descent functions and the
one we get using the most optimal theta value, 'theta_best'*

```
pred_best = X.dot(theta_best)
print("Error from Optimal Theta : {}".format(mean_squared_error(Y,
                                                                pred_best)))

pred_gd = X.dot(theta_gd)
print("Error from Calculated Theta for Batch GD: {}".format(
    mean_squared_error(Y, pred_gd)))

pred_mbgd = X.dot(theta_mbgd)
print("Error from Calculated Theta for Mini-Batch GD: {}".format(
    mean_squared_error(Y, pred_mbgd)))

pred_sgd = X.dot(theta_sgd)
print("Error from Calculated Theta for Stochastic GD: {}".format(
    mean_squared_error(Y, pred_sgd)))

print("Error from Calculated Theta for sklearn SGD: {}".format(
    mean_squared_error(Y, pred_sklearnSGD)))
```

```
Error from Optimal Theta : 21.894831181729206
Error from Calculated Theta for Batch GD: 24.698483543800545
Error from Calculated Theta for Mini-Batch GD: 21.96812758471529
Error from Calculated Theta for Stochastic GD: 22.845951445106508
Error from Calculated Theta for sklearn SGD: 22.82971463980024
```


```
In [49]: print("Comparing the weights from our SGD and sklearn Regressor: ")
sklSGD = clf.coef_
for i in range(len(sklSGD)):
    print(" {0:2d} : {1: 10.8f} - {2: 10.8f}".format(
        i+1,
        theta_sgd[i][0],
        sklSGD[i]))
```

```
Comparing the weights from our SGD and sklearn Regressor:
1 : 22.46789267 - 11.31434460
```

2 :	-0.66849967	-	-0.66531255
3 :	1.14203826	-	0.48237328
4 :	0.29153106	-	-0.39389998
5 :	0.64938312	-	0.81531448
6 :	-1.95053319	-	-1.06691459
7 :	2.82412889	-	3.25478057
8 :	-0.28268455	-	-0.14535041
9 :	-2.94866511	-	-2.10835523
10 :	2.86720407	-	0.86533220
11 :	-2.06102117	-	-0.39270045
12 :	-1.61062724	-	-1.82439325
13 :	1.28868717	-	0.93486820
14 :	-4.43665434	-	-3.27933427

Comparing the predictions of our SGD, sklearn SGD and the actual Values

```
In [58]: # Take 25 samples from the 506 samples
x_axis = [i for i in range(samples) if i%20==0]
plt.figure(figsize=(10.0, 8.0))
plt.scatter(x_axis, Y[x_axis], marker='x',
            color='r', label='Actual Values')
plt.scatter(x_axis, pred_sklearnSGD[x_axis], marker='o',
            color='b', label='Sklearn SGD')
plt.scatter(x_axis, pred_sgd[x_axis], marker='d',
            color='g', label='Our SGD')
plt.xlabel("Samples")
plt.ylabel("Predicted Values")
plt.title("Comparison of Predicted Prices")
plt.legend()
plt.show()
```



In [59]: *# Hence, we see we were able to beat the SGD regressor score but we couldn't
surpass the score we got from the optimal value of theta*

Conclusions

- Batch Gradient descent is the fastest to execute, followed by mini-batch gradient descent. The stochastic gradient descent is immensely slow by around 95 times compared to batch.
- On the other hand, we see stochastic gradient descent is fastest to converge around the optimal theta and batch gradient descent is the slowest. This might be an observation specific to this dataset. In a linear regression dataset a few randomly sampled points can give a good approximation of the optimal plane(which we are trying to fit) and that helps the stochastic and mini-batch to converge faster. Whereas, the batch gradient while trying to fit the theta plane over all the points in each iteration, becomes slow. In more complex datasets the stochastic and mini-batch are expected to show a much more random behaviour.
- The mini-batch gradient descent came closest to the optimum values of theta,

which we calculated directly by solving the vectorised normal equation. The mean squared error in this case was even better than the sklearn SGDRegressor error value. The stochastic gradient descent also came close, but the batch gradient descent has little amount of error.

- Overall, we can say considering the time and optimal value of theta, that the mini-batch gradient descent with batch size of 25 showed the best performance.

In []: